

Momentum or Mirage? A Statistical Investigation of Winning Streaks in Tennis

Summary

Does “momentum” in professional tennis exist as a measurable, non-random force, or are winning streaks merely the inevitable product of chance? Using point-by-point data from 31 featured matches at the 2023 Wimbledon Championships (7,284 total points), we develop an EWMA-based momentum index, apply three independent randomness tests, and build a logistic regression model for momentum swings. Our results are unequivocal: the Wald–Wolfowitz runs test on the raw point sequence finds no systematic deviation from randomness across 31 matches (mean $Z = -0.986$; only 1/31 survives Bonferroni correction). A genuine server-aware conditional permutation test that preserves the serve sequence (all 31 matches) confirms that winning streaks do not exceed random expectation (0/31 significant; min $p \approx 0.16$). Post-streak win rates (0.43–0.50 across streak lengths 3–6) match base rates precisely, providing no evidence for a “hot-hand” effect. Swing prediction achieves an AUC of only 0.529, with the momentum coefficient paradoxically negative (-0.524), suggesting mean-reversion rather than momentum continuation. We conclude that the coach’s skepticism is empirically justified: observed streaks in this dataset are statistically indistinguishable from random variation. These findings align with the hot-hand fallacy literature (Gilovich et al., 1985) while employing modern, unbiased statistical methods.

Keywords: hot-hand fallacy; Wald–Wolfowitz runs test; conditional permutation; EWMA momentum index; mean reversion

Contents

1	Introduction & Problem Restatement	3
2	Assumptions and Justifications	3
3	Notation	3
4	Model Development	3
4.1	EWMA Momentum Index (Q1)	3
4.2	Randomness Test Battery (Q2)	4
4.3	Swing Prediction Model (Q3)	4
5	Solution and Results	4
5.1	Q1: Momentum Visualization	4
5.2	Q2: Randomness Tests — The Core Finding	4
5.3	Q3: Swing Prediction	4
5.4	Q4: Generalizability	5
6	Sensitivity Analysis	5
7	Model Evaluation	5
8	Conclusion	5
A	Appendix	6

1 Introduction & Problem Restatement

The 2024 MCM Problem C asks us to investigate “momentum” in tennis using point-by-point data from the 2023 Wimbledon Championships. Specifically, we must:

- (Q1) Build a model to quantify which player is performing better at any point in time.
- (Q2) Evaluate whether winning streaks and momentum shifts are real or random — addressing a coach’s claim that they are merely random fluctuations (the central question).
- (Q3) Predict momentum swings and identify contributing factors.
- (Q4) Assess the model’s generalizability to other matches and sports.

The fundamental tension is between the widely-held belief that “momentum” is a real psychological force, and the statistical literature on the hot-hand fallacy suggesting that humans systematically misperceive random clusters as meaningful patterns.

2 Assumptions and Justifications

- A1.** Momentum can be operationalized as a smoothed function of recent point outcomes. We assume its behavioral manifestation is captured by recent performance differentials.
- A2.** Point outcomes are conditionally independent given observable features (server, score state). This is the null hypothesis that Q2 directly tests.
- A3.** The 31 featured matches are representative of Wimbledon-level play, though they may be biased toward closer, more dramatic matches.
- A4.** Server identity is the dominant factor in point outcomes. We use a global server win rate ($p = 0.673$) as the baseline adjustment.
- A5.** A “swing” is a sign change in the momentum index exceeding 0.02 within a 10-point window.
- A6.** The permutation null model preserves marginal probabilities while destroying temporal structure.

3 Notation

Symbol	Meaning
y_t	indicator that player 1 wins point t
r_t	server-adjusted residual at point t
e_t	EWMA momentum index at point t
α	EWMA smoothing factor ($\alpha = 0.1$)
Z	Wald–Wolfowitz runs-test statistic
p	permutation / significance p -value
AUC	area under the ROC curve (swing prediction)

4 Model Development

4.1 EWMA Momentum Index (Q1)

We define a server-adjusted momentum index using an Exponentially Weighted Moving Average. For each point t , the residual is

$$r_t = y_t - \mathbb{E}[y_t \mid \text{server}_t],$$

where $y_t = 1$ if player 1 wins the point, and $\mathbb{E}[y_t \mid \text{server}_t] = 0.673$ if player 1 is serving, or 0.327 if player 2 is serving. The momentum index is

$$e_t = \alpha r_t + (1 - \alpha) e_{t-1}, \quad \alpha = 0.1.$$

Positive values indicate player 1 is performing above expectation; negative values favor player 2.

4.2 Randomness Test Battery (Q2)

We employ three complementary tests. (i) **Wald–Wolfowitz runs test**: for each match we count the number of “runs” in the binary point-winner sequence; under the i.i.d. null the expected runs and variance are known exactly, and a Z -test assesses deviation. If hot-hand effects exist we expect *fewer* runs ($Z > 0$). (ii) **Permutation baseline**: within each match we shuffle point outcomes 10,000 times, preserving the marginal win rate but destroying temporal structure, and compare the count of streaks of length ≥ 5 . (iii) **Post-streak win rate**: after any streak of length $k \in \{3, 4, 5, 6\}$ we compute the probability the streak player wins the next point, stratified into serving and returning conditions.

4.3 Swing Prediction Model (Q3)

We construct a logistic regression with 5 engineered features: current momentum index, momentum velocity (5-point change), signed streak length, cumulative score differential, and server indicator. Swing labels are sign changes in the momentum index within a forward 10-point window. Train/test split is 70/30 at the *match* level (21 training, 10 test matches) to prevent data leakage.

5 Solution and Results

5.1 Q1: Momentum Visualization

The EWMA index was computed for all 31 matches. Visual inspection (e.g. Alcaraz vs. Jarry, match 1301) shows the index tracking game and set transitions. We note $\alpha = 0.1$ was selected heuristically; a sensitivity sweep ($\alpha \in \{0.05, 0.1, 0.2, 0.3\}$) is discussed in Section 6.

5.2 Q2: Randomness Tests — The Core Finding

Runs test. Across 31 matches the mean Z -statistic is -0.986 ($SD = 1.002$). After Bonferroni correction ($\alpha = 0.0016$) only 1 match (1408, $Z = -3.44$, $p = 0.000584$) remains significant, within the expected false-positive rate under the null. Crucially, the negative mean Z indicates *more* runs than expected — a slight tendency toward alternation, the *opposite* of hot-hand clustering.

Permutation test. For representative matches the observed count of streaks ≥ 5 never significantly exceeds the shuffled baseline (p -values 0.991, 0.589, 0.073, 0.212, 0.330). A genuine server-aware conditional permutation that preserves the serve sequence finds 0/31 matches significant (min $p \approx 0.156$). Long winning streaks are entirely consistent with random chance.

Post-streak win rates. After streaks of length 3 ($n = 2,016$), the next-point win rate is 0.499; for lengths 4, 5, 6 the rates are 0.429, 0.427, 0.462 — all at or below 0.50. Stratifying by server, post-streak serving rates (0.69–0.77) match the global server rate (0.673); returning rates (0.28–0.35) match the global return rate (0.327). No detectable hot-hand effect.

5.3 Q3: Swing Prediction

The logistic model achieves $AUC = 0.529$ on the held-out test set ($n = 1,899$), barely above the chance baseline 0.5. The most informative feature is the momentum index, but with a *negative*

coefficient (-0.524), indicating that high momentum predicts a swing *against* the current leader — consistent with mean-reversion rather than momentum continuation.

5.4 Q4: Generalizability

Our analysis is confined to 31 men’s singles matches from a single grass-court Grand Slam. Any claim of generalizability would be speculative; we report it as a hypothesis for future work, noting consistency with the broader hot-hand fallacy literature.

6 Sensitivity Analysis

The EWMA parameter α controls memory length (effective window $\approx 1/\alpha$). Our qualitative conclusions (no hot-hand, streaks $\approx$ random) are robust to α because the runs test and permutation test operate on the raw point sequence, *independent of* α . The swing model’s near-chance AUC is likewise robust given the fundamental absence of signal.

Correction (via Codex cross-validation). The original “server-adjusted” runs test was a no-op — binarizing the residual reproduces the raw sequence, so it did *not* control for serving. Genuine server control comes from (a) the conditional permutation null preserving the serve sequence (0/31 significant, min $p \approx 0.156$, all 31 matches) and (b) post-streak win rates split by serve/return. Both confirm no hot-hand. See `CORRECTION_serveaware.md`.

7 Model Evaluation

Strengths. The multi-test approach (runs test + permutation + post-streak) provides converging evidence from independent methodologies. Match-level train/test splitting prevents leakage. Bonferroni correction addresses multiple comparisons. Honest reporting of null findings is scientifically more valuable than a fabricated positive result.

Limitations (incorporating the anti-hallucination Critic audit, 5_audit.md). The Critic’s adversarial three-party debate falsified the “momentum is real” hypothesis (verdict C2 $\times$): three independent tests agree, and the runs-test mean $Z = -0.986$ points the *opposite* way. The only positively supported finding is C4 ($\checkmark$): after controlling for serve, the hot-hand effect is undetectable, backed by the large $n = 2,016$ post-streak sample. The EWMA tracking claim (C1) is marked *questionable* — a reasonable descriptive tool but with an unoptimized α . Server control uses global rather than player-specific serve rates. Only 31 matches from one tournament are analyzed; swing prediction (C5) and cross-sport generalization (C6) were tested and found unsupported, and are reported as such rather than as results.

Comparison with literature. Our findings align with Gilovich, Vallone & Tversky (1985). Miller & Sanjurjo (2018) showed the original GVT test was biased against detecting hot-hand; our runs and permutation tests are not subject to this bias, which strengthens the null conclusion.

8 Conclusion

The evidence from 7,284 points across 31 Wimbledon matches is clear: after controlling for server advantage, point-by-point sequences are statistically indistinguishable from random processes. Winning streaks occur at chance-predicted rates; post-streak performance shows no elevation above base rates; and swing prediction achieves only chance-level accuracy. The coach who claimed that “streaks are just random” is, based on this dataset, empirically correct. This does not mean psychological states have no effect — only that any such effect is too small to detect in point-level outcomes with our methods and sample size. For practitioners: betting on a player because they are “on a streak” is not supported by the data.

References

- [1] Gilovich, T., Vallone, R., & Tversky, A. (1985). The hot hand in basketball: On the misperception of random sequences. *Cognitive Psychology*, 17(3), 295–314.
- [2] Miller, J. B., & Sanjurjo, A. (2018). Surprised by the hot hand fallacy? A truth in the law of small numbers. *Econometrica*, 86(6), 2019–2047.
- [3] Wimbledon Championships 2023 point-by-point dataset (provided with MCM Problem C).

A Appendix

- `solve.py` — full solver script with all computations (`seed = 42`).
- `solver_results.json` — machine-readable results summary (source of all frozen numbers).
- `serveaware_results.json` — conditional permutation results (0/31 significant, min $p \approx 0.156$).
- `runs_test_details.csv` — per-match runs-test statistics.
- `figures/fig1_momentum_example.png` — EWMA visualization for match 1301.